



Une école de l'IMT

Convolutional and Deep Neural Networks

A. Fiandrotti

TSIA203

Introduction to Deep Learning

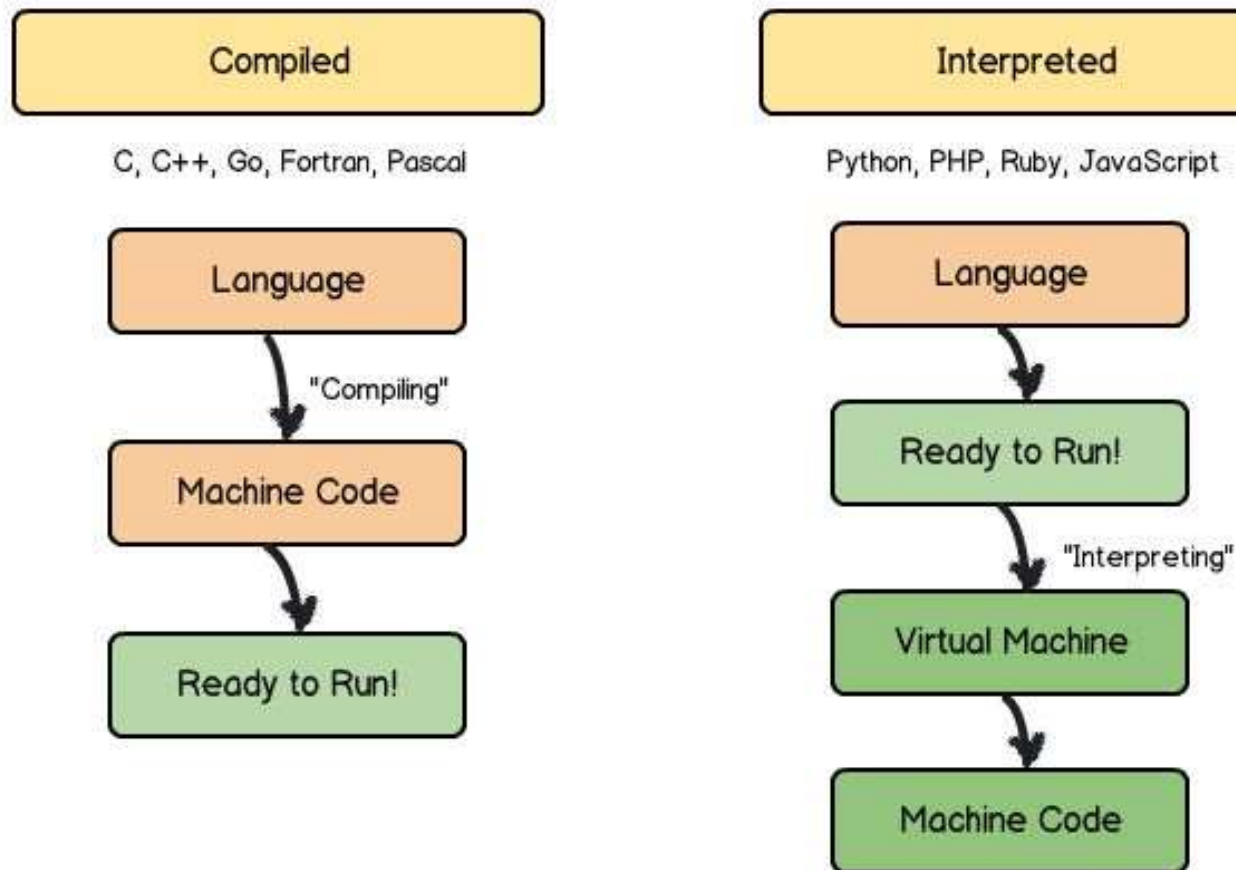




Part 4

Introduction to Python

Compiled vs Interpreted Languages





Python Versions and Syntax

■ Current version 3.x

```
print ("Hello World!")
```

■ Old 2.x code not 100% compatible

```
print "Hello World!"
```

■ Always make sure to use 3.x APIs!

- PyTorch, Keras, TF follow Python 3.0



Indentation

■ Python requires indentation

- Tabulation, white spaces

```
if myIntVar == 10:
    print ("First branch taken!")
    if myStrVar == "Some text"
        print ("Second branch taken!")
```

■ Indentation error example

```
>>> if myVar == 10:
...   print ("Branch taken!")
      File "<stdin>", line 2
        print ("Branch taken!")
          ^
IndentationError: expected an indented block
```



Indentation

■ Python requires indentation

```
if myIntVar == 10:
    print ("First branch taken!")
    if myStrVar == "Some text"
        print ("Second branch taken!")
```

■ Indentation error example

```
>>> if myVar == 10:
... print ("Branch taken!")
File "<stdin>", line 2
    print ("Branch taken!")
    ^
IndentationError: expected an indented block
```

Variables Types

■ Weakly typed variables

```
# boolean variable
```

```
>>> v=True
```

```
>>> type(v)
```

```
<type 'bool'>
```

```
# integer number
```

```
>>> v=1
```

```
>>> type(v)
```

```
<type 'int'>
```

```
# floating point number
```

```
>>> v=1.0
```

```
>>> type(v)
```

```
<type 'float'>
```

```
# text string
```

```
>>> v="Hello World!"
```

```
>>> type(v)
```

```
<type 'str'>
```

```
# defining a list
```

```
>>> list=['Rat','Cat','Dog']
```

```
>>> type(list)
```

```
<type 'list'>
```

```
# defining a list of lists
```

```
>>> matrix=[[1,2,3],[4,5,6]]
```

```
>>> type(matrix)
```

```
<type 'list'>
```

```
>>> len(matrix)
```

```
2
```

```
>>> len(matrix[0])
```

```
3
```

```
# defining a dictionary
```

```
>>> dictionary={'Rat', 10, True}
```

```
>>> type(dictionary)
```

```
<type 'dict'>
```

Variables by Reference

```
# creating list
>>> m=[1, 2, 3]

# creating second list pointer
>>> m2 = m

# modifying list via first pointer
>>> m[0][1] = -1
>>> m2
[-1, 2, 3]

# dropping first pointer
>>> del(m)
>>> type(m)
Traceback (most recent call last):
NameError: name 'm' is not defined

# list still accessible via second pointer
>>> type(m2)
<class 'list'>
```




Variables Cloning

■ See also *deepcopy()*

```
# original list
>>> list = [0, 1, 2]

# cloning the list
>>> new_list = list.copy()

# appending one element to the cloned list
>>> new_list.append(4)

# printing new and old list
>>> print('Old List: ', list)
Old List: [0, 1, 2]
>>> print('New List: ', new_list)
New List: [0, 1, 2, 4]
```



Conditional Branches

■ Mind the indentation!

```
>>> v=1
>>> if v == 0:
...     print("variable equal to 0")
... elif v == 1:
...     print("variable equal to 1")
... else:
...     print("variable value is " + str(v))
...
variable equal to 1
```



For Loops

■ The C / Java way

```
# define a list
>>> list=['a',-1,1.0]

# getting length of list
>>> length = len(list)

# Iterating the index
>>> for i in range(length):
... print(list[i])
a
-1
1.0

# same as 'for i in range(len(list))'
```



For Loops

■ The Python way

```
# define a list
>>> list=['a',-1,1.0]

# iterate through the list elements the python way
>>> for e in list:
...     print ("element " +str(e) + " is of type " + str(type(e)))
...
element a is of type <class 'str'>
element -1 is of type <class 'int'>
element 1.0 is of type <class 'float'>
```



While Loops

■ Also do ... while

```
# creating counter
>>> cnt = 5

# cycling throuh
>>> while cnt > 0:
...     print (cnt)
...     cnt = cnt - 1
...
5
4
3
2
1
```



Functions

- Positional arguments
- Keyword arguments

```
>>> def f(x, y=1, z=1):  
...     return x + y + z  
...  
>>> print(f(1, 2, 3))  
6  
>>> print(f(1))  
3  
>>> print(f(1, 2))  
4  
>>> print(f(1, z=3))  
5  
>>> print(f(1, z=3, y=2))  
6  
>>> print(f(z=3, 1))  
File "<stdin>", line 1  
SyntaxError: positional argument follows keyword argument
```



Importing modules

- Importing entire libraries or subcomponents

```
>>> import numpy
>>> numpy.array([1, 2])
array([1, 2])
>>> import numpy as np
>>> np.array([1, 2])
array([1, 2])
>>> from numpy import array
>>> array([1, 2])
array([1, 2])
>>> from numpy import array as ar
>>> ar([1, 2])
array([1, 2])
```



Python Versions and Syntax

- **Current version 3.x (PyTorch, Keras, TF, ...)**

```
print ("Hello World!")
```

- **Old 2.x code not 100% compatible**

```
print "Hello World!"
```

- **How do I know my python version / path ?**

```
attilio@debian:~$ python --version  
Python 3.5.3
```

```
attilio@debian:~$ which python  
/usr/bin/python
```




The Anaconda Distribution

- **Problem: my OS has python 2.x, but I need 3.x**
 - Python 2.x needed by OS, have no root rights, ...
- **Solution: install [Anaconda](#)**
 - No need to be root (installed in user's home)

```
attilio@debian:~$ which python
/home/attilio/anaconda3/bin/python
```

- **Manage libraries via *conda***

\$conda search [library name]

\$conda install [library name]

Conda Package Manager

```
attilio@debian:~$ conda -V  
conda 3.7.0
```

```
attilio@debian:~$ conda search keras  
Loading channels: done
```

# Name	Version	Build	Channel
keras	1.1.1	py27_0	pkgs/free
keras	1.1.1	py34_0	pkgs/free
keras	1.1.1	py35_0	pkgs/free
keras	2.1.6	py27_0	pkgs/main
keras	2.1.6	py35_0	pkgs/main
keras	2.1.6	py36_0	pkgs/main

```
attilio@debian:~$ conda install [-c XYZ] keras[=2.1.6]=py27_0]
```



Python Environments

■ Create the environment

```
attilio@debian:~$ conda create -n yourenvname python=x.x anaconda
```

■ Activate the environment

```
attilio@debian:~$ source activate yourenvname
```

■ Install packages in the environment

```
attilio@debian:~$ conda install -n yourenvname [package]
```

■ Deactivate the environment

```
attilio@debian:~$ source deactivate
```

■ Delete the environment

```
attilio@debian:~$ conda remove -n yourenvname -all
```



Part 5

Introduction to Keras

PyTorch



- ▶ *Torch7* tensor library developed at EPFL (C++)
- ▶ *Torch* was a *Lua* binding to *Torch7*



ÉCOLE POLYTECHNIQUE
FÉDÉRALE DE LAUSANNE



The Keras Language

- ▶ High level framework for machine learning
 - ▶ *High-level* w.r.t. PyTorch, TensorFlow, etc.
- ▶ Several backends available
 - ▶ We will use the TensorFlow backend (Google)

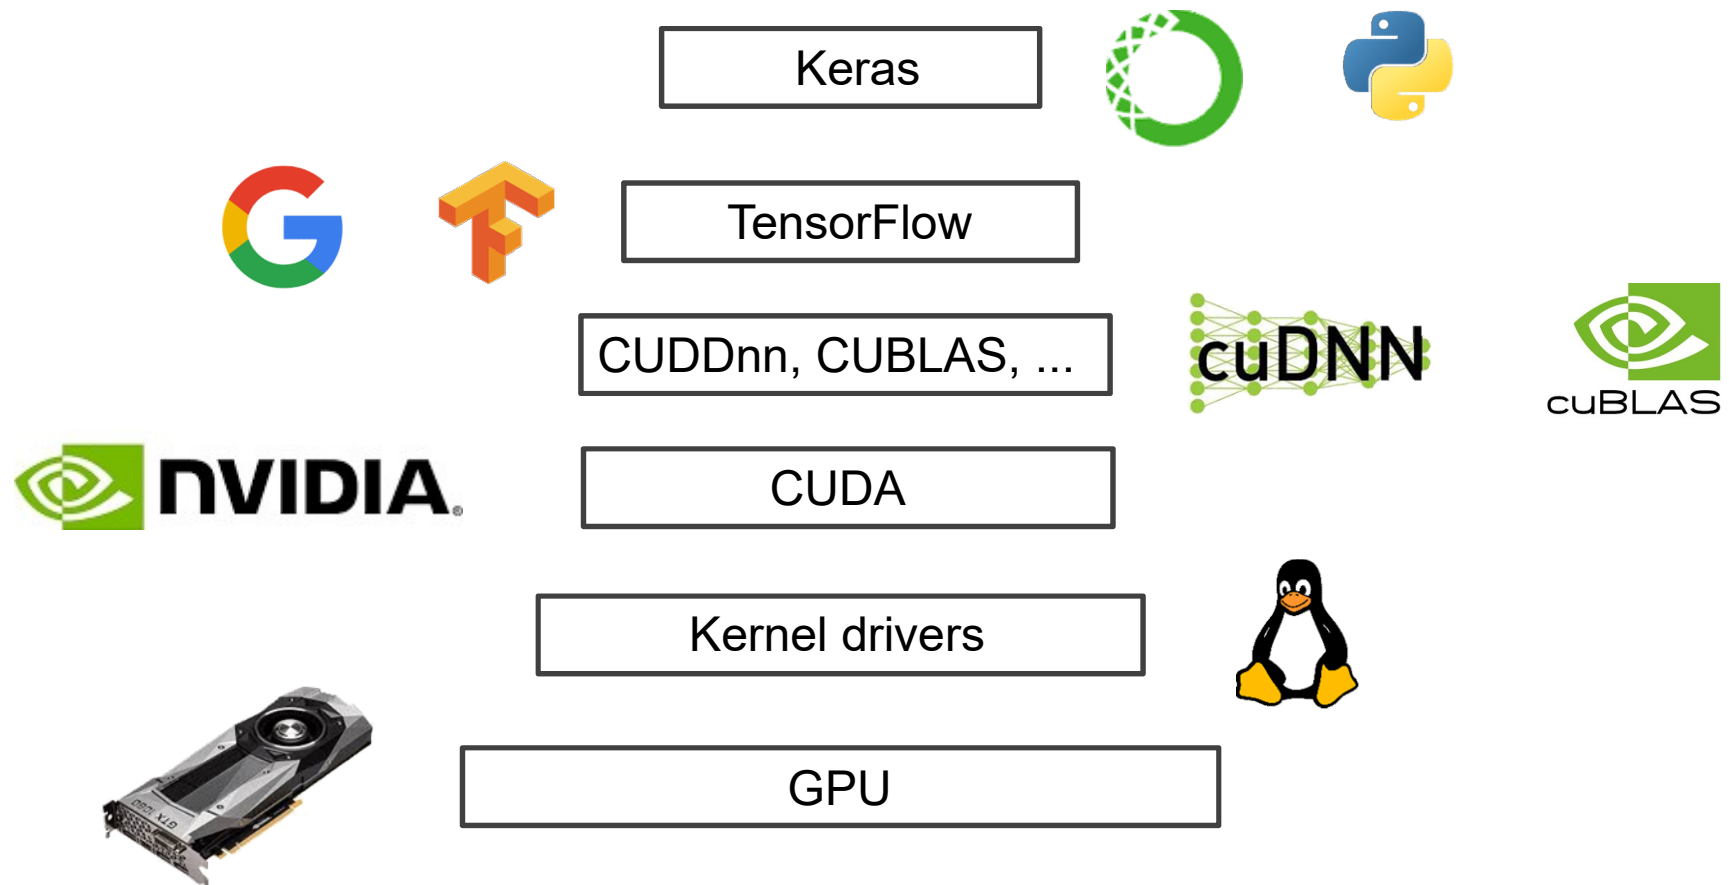


The Keras Language

- ▶ High level framework for machine learning
 - ▶ *High-level* w.r.t. PyTorch, TensorFlow, etc.
- ▶ Several backends available
 - ▶ We will use the TensorFlow backend (Google)
- ▶ Main author François Chollet
 - ▶ Google employee, former ENSTA Paristech alumni



Keras System Stack

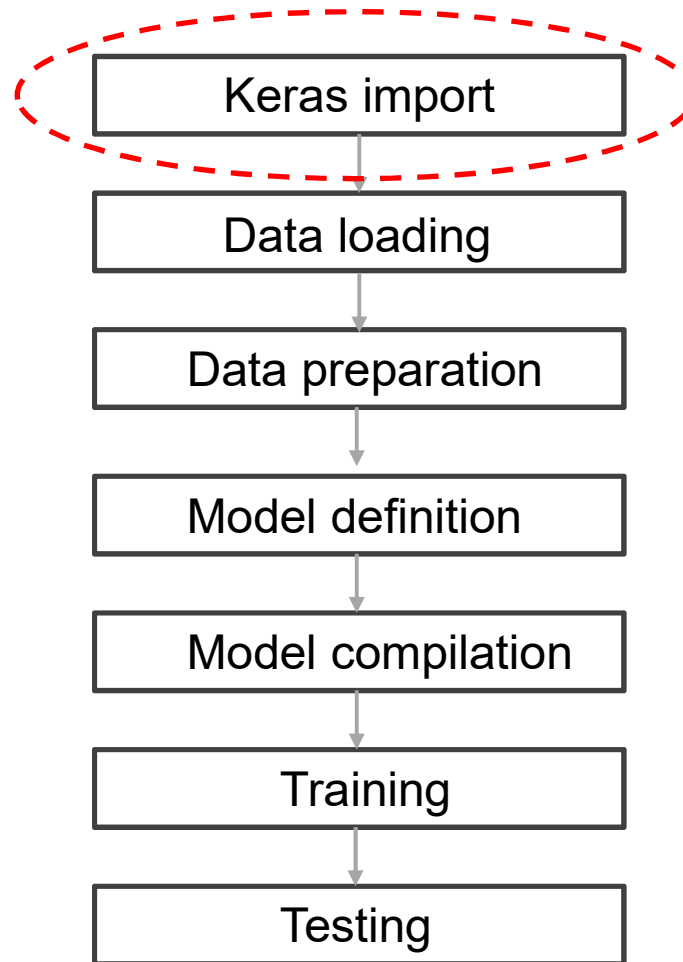


Keras



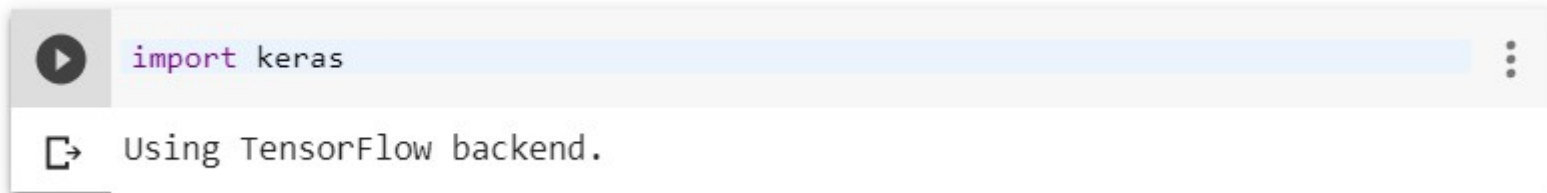
- ▶ Datasets loading
 - ▶ Popular datasets such as MNIST, etc available
 - ▶ Uses *scikit-learn* for synthetic data
- ▶ Defining network architectures
 - ▶ Non-sequential models supported
 - ▶ Pretrained deep models (AlexNet, ResNet)
- ▶ Training a network
 - ▶ Multiple optimizers available
 - ▶ One-line *fit()* function
- ▶ Visualizing data and results
 - ▶ *Relies on matplotlib* for visualization

Typical Keras Dataflow



Module Import – Backend Ordering

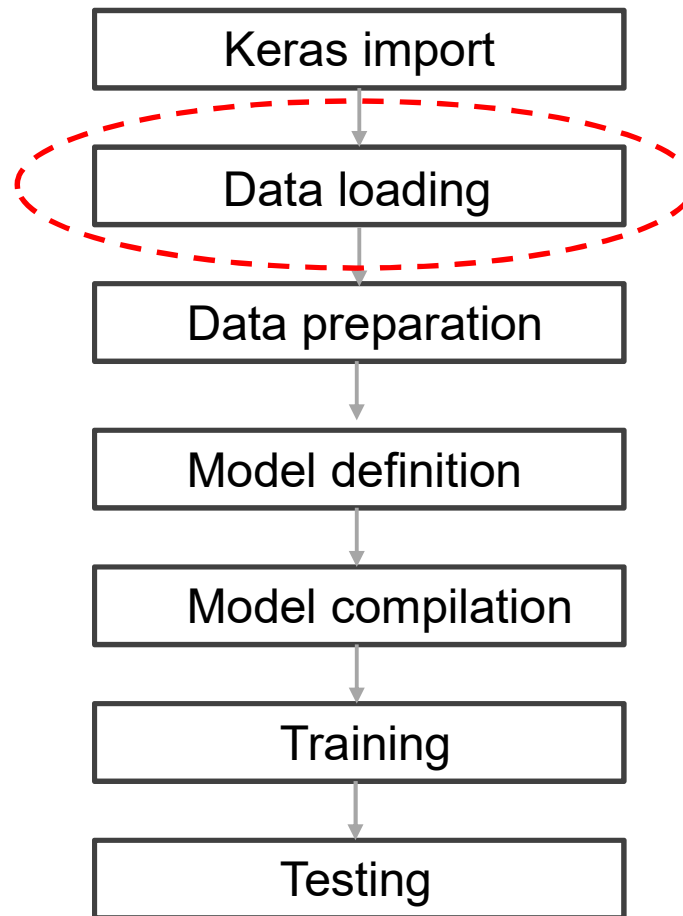
- ▶ Make sure Keras uses the *TensorFlow* backend
- ▶ «*NHWC*» data ordering required for images
 - ▶ *N* -> *image index in batch*
 - ▶ *H* -> *image height*
 - ▶ *W* -> *image width*
 - ▶ *C* -> *image channel*



```
import keras
```

Using TensorFlow backend.

Typical Keras Dataflow




Data Generation

- ▶ Will use the *sklearn* beackend

```
from sklearn import datasets
```

- ▶ Generate two Gaussian clusters of points

```
data, labels = datasets.make_classification(  
    n_samples=100,  
    n_features=2, n_informative=2, n_redundant=0,  
    n_classes=2, n_clusters_per_class=1,  
    class_sep=1.5, flip_y=0,  
    shuffle=True)
```



Data Generation

- ▶ Will use the *sklearn* beackend

```
from sklearn import datasets, model_selection
```

- ▶ Generate two Gaussian clusters of points

```
data, labels = datasets.make_classification(  
    n_samples=100,  
    n_features=2, n_informative=2, n_redundant=0,  
    n_classes=2, n_clusters_per_class=1,  
    class_sep=1.5, flip_y=0,  
    shuffle=True)
```

- ▶ Separate 90% train and 90% test

```
train_data, test_data,  
train_labels, test_labels = model_selection.train_test_split(  
    data, labels,  
    test_size = 0.1)
```

Data Generation - Visualization


► Matlab-like plotting library

```
import matplotlib.pyplot as plt
```

► Plot the data

```
colors = ['steelblue' if label == 1 else 'darkred' for label in labels]
plt.scatter(X[:,0], X[:,1], color=colors)
plt.show()
```

```
Y.shape, X.shape
((100,), (100, 2))
```



Dataset Loading - MNIST

- ▶ Keras has some popular datasets pre-packaged

```
from keras.datasets import mnist
```

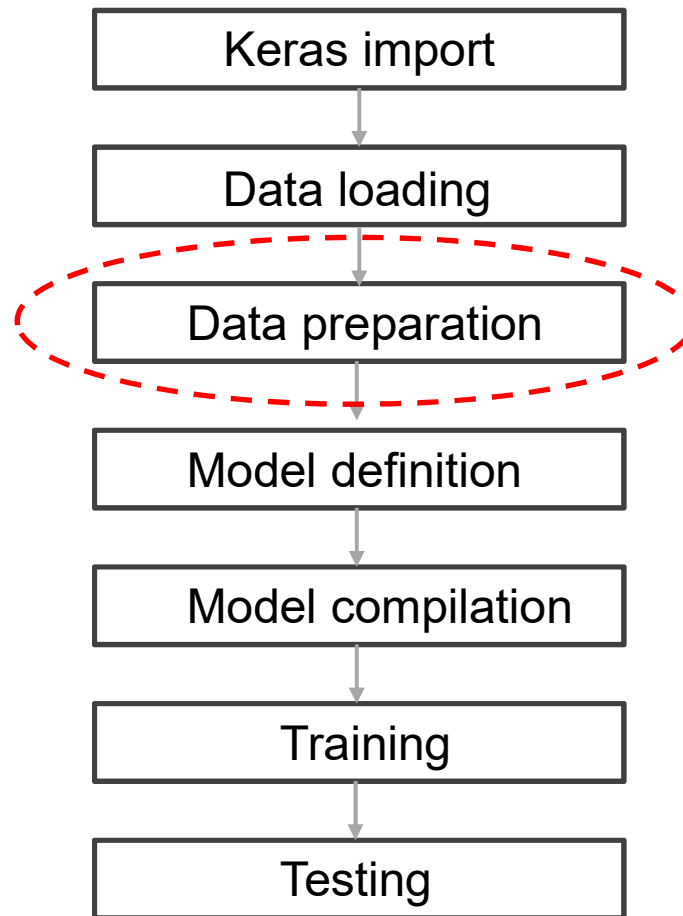
- ▶ Load images and labels into memory

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()

print(y_train.shape)
(60000,)
print(x_train.shape)
(60000, 28, 28)
```

- ▶ Train samples are a `numpy.ndarray` of `int8`
- ▶ `N = 60000` samples
 - ▶ `W x H = 28 x 28` px. samples grayscale

Typical Keras Dataflow



Data Preparation - Images

- ▶ Load images and labels into memory

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()

print(x_train.shape)
(60000, 28, 28)
```

- ▶ Samples are in (N,W,H) ordering
 - ▶ The TF backend requires (N,W,H,C) ordering
- ▶ Samples are 256-grayscale `int8` arrays
 - ▶ Neural networks require `float` in input
- ▶ Samples are in 0-255 interval
 - ▶ Normalization desirable

Data Preparation - Images

- ▶ Let us exploit **numpy** builtins!
- ▶ **Reshape** the samples to NWHC order

```
x_train = numpy.reshape(x_train, newshape, order='C')
```

- ▶ **Recast** the samples from **int8** to **float**

```
x_train = x_train.astype('float_32')
```

- ▶ **Squeeze** from [0-255] to [0-1]

```
x_train = x_train /= 255
```

- ▶ **Normalize** to have zero-**mean** (and unit-**std**)

```
x_train = (x_train - x_train.mean())  
          -----  
          x_train.std()
```

- ▶ What you do on train samples, do it also on test samples
 - ▶ Yet on train statistics!

Data Preparation - Images



- ▶ Labels (classes) are encoded as integers [0-9]

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()

print(y_train.shape)
(60000,)
```

- ▶ One-hot encoding required (multiclass problem)
 - ▶ Use `to_categorical()`

```
from keras.utils.np_utils import to_categorical
y_train_oh = to_categorical(y_train)

print(y_train_oh.shape)
(60000, 10, 2)

print(train_labels.shape)
[[1. 0.] [1. 0.] [1. 0.] [1. 0.] [1. 0.] [0. 1.] [1. 0.] [1. 0.] [1. 0.] [1. 0.]]
```



Data Loading - Generating Data

- ▶ Transform your labels to *one-hot* encoding first

```
from keras.utils.np_utils import to_categorical

data, labels = datasets.make_classification(
    n_samples=5, n_features=2, n_classes=2, [...])

labels_oh = to_categorical(labels)
```

```
print(data)
array([[ 3.34195848, -2.02906319],
       [ 0.12096805,  2.27640523],
       [ 2.02600963, -0.62195723],
       [ 1.91840064, -2.19699255],
       [ 3.47918424, -0.43343625]])
```

```
print(labels)
array([0,
       1,
       0,
       0,
       1])
```

```
print(labels_oh)
array([[1, 0],
       [0, 1],
       [1, 0],
       [1, 0],
       [0, 1]])
```

Data Preparation - Augmentation

- ▶ Import ImageDataGenerator

```
from keras.preprocessing.image import ImageDataGenerator
```

- ▶ Create proper ImageDataGenerator instance

```
myDatagen = ImageDataGenerator(  
    rotation_range=0,  
    width_shift_range=0.1,  
    height_shift_range=0.1,  
    horizontal_flip=True,  
    vertical_flip=False  
)
```

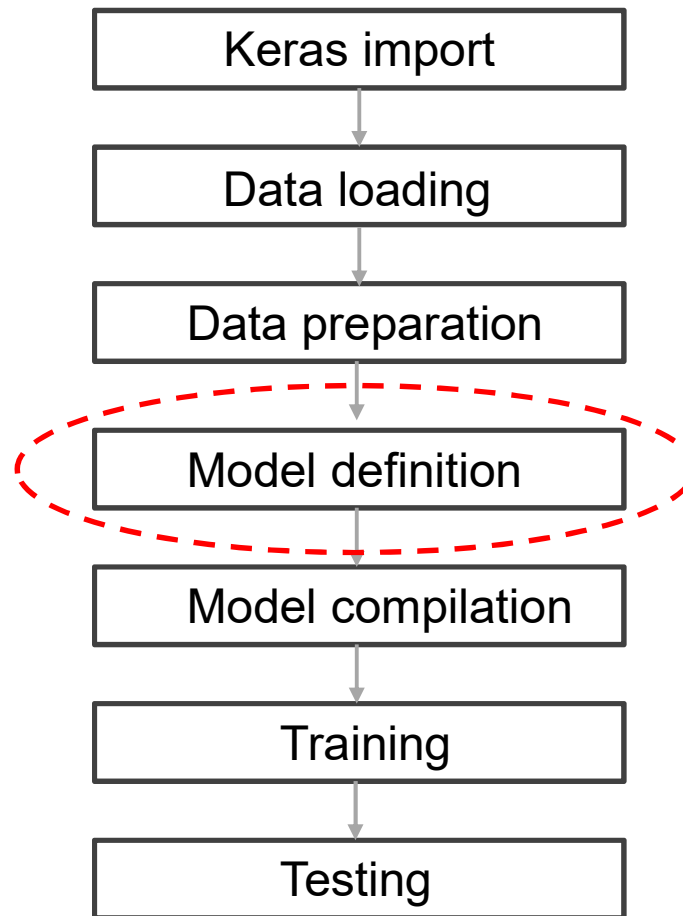
- ▶ Can optionally normalize data for you

- ▶ Must fit generator to your data if normalization used!

- ▶ Further steps required during training

- ▶ Detailed later

Typical Keras Dataflow



Model definition – Sigmoid Classifier


► Import keras modules

```
from keras.models import Sequential
from keras.layers import Dense, Activation
```

► Use the *sequential API*

```
model = Sequential()
model.add(Dense(1, input_dim=2))
model.add(Activation('sigmoid'))
```

```
model = Sequential(
    [Dense(32, input_dim=2),
     Activation('sigmoid'),])
```



► Use *functional API* for complex models

► forks, yields, loops, etc.

Model definition – Convolutional Layers

- ▶ Import keras modules

```
from keras.models import Conv2D, MaxPooling2D
```

- ▶ Convolutional layer with ReLU activation and MaxPooling

```
model = Sequential()  
model.add(Conv2D(filters=6,  
                 kernel_size=(5,5),  
                 input_shape=(w,h,c),  
                 padding='[valid|same]',  
                 data_format='channels_last'))  
model.add(Activation('relu'))  
Model.add(MaxPooling2D(pool_size=2))
```

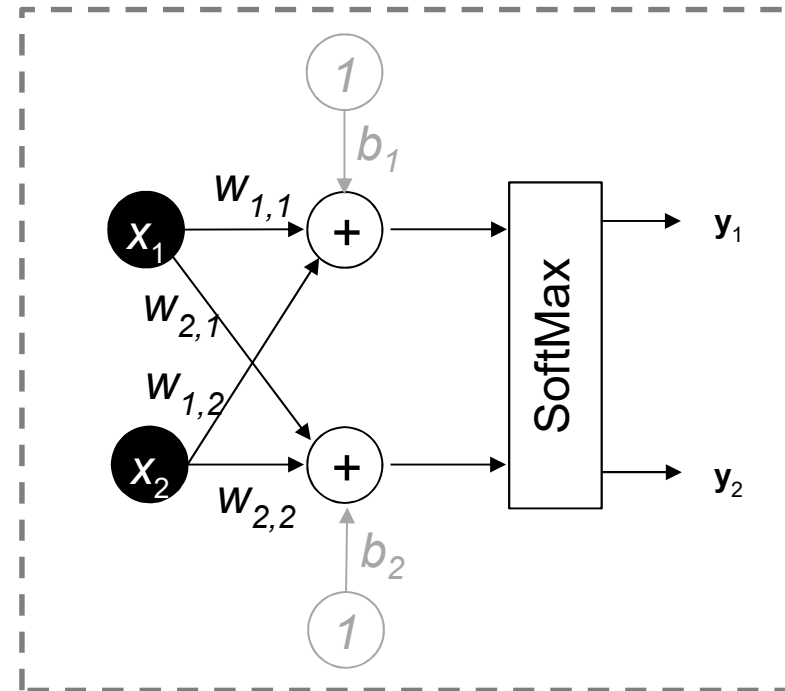
- ▶ Serialize the feature maps into feature vectors

```
model.add(Flatten())
```

Model definition - Multiple Classes Output

- ▶ Need to adjust
 - ▶ Number of nodes in output layer
 - ▶ Replace Sigmoid activation with Softmax

```
model = Sequential()  
model.add(Dense(2, input_dim=2))  
model.add(Activation('softmax'))
```



Model definition - Regularizers

- ▶ Implemented as layer parameters

```
from keras import regularizers
```

- ▶ L1 and L2 norm regularizers commonly used
- ▶ Apply individually to each layer

```
model = Sequential()

model.add(Dense(1000,
                kernel_regularizer=regularizers.l2(0.01),
                bias_regularizer=regularizers.l2(0.01)))

model.add(Activation(...))
model.add(Dense(10, ...))
model.add(Activation('softmax'))
```

Model definition - Dropout

- ▶ Implemented as a layer

```
from keras.models import Dropout
keras.layers.Dropout(rate, noise_shape=None, seed=None)
```

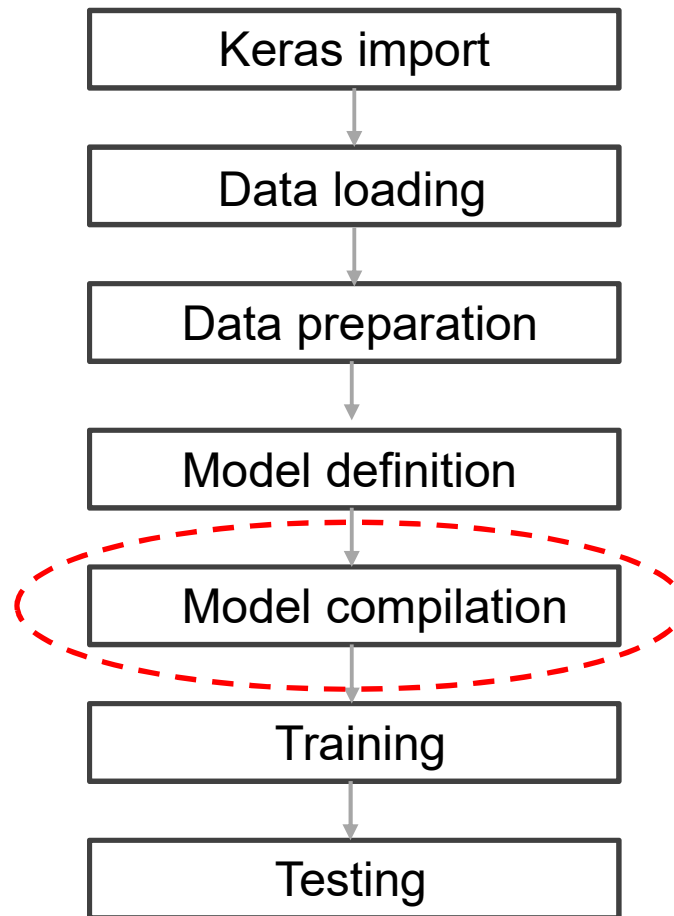
- ▶ Usually useful before most parametrized layer

```
model = Sequential()
model.add(Conv2D(...))
model.add(Activation(...))
model.add(MaxPooling2D(...))

model.add(Dropout(rate = 0.5))

model.add(Dense(1000))
model.add(Activation(...))
model.add(Dense(10, ...))
model.add(Activation('softmax'))
```

Typical Keras Dataflow



Model compilation – Binary Classification

► Define an optimizer

```
myOpt = SGD(lr=0,05, decay=10e-6)
```

► Define loss (and performance) metric

```
model.compile(optimizer=myOpt, loss='mse', metrics=['accuracy'])
```

► After compiling the model, visualize it

```
model.summary()
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 1)	3
Total params: 3		

Model compilation - Multiple classes

► Define an optimizer

```
myOpt = SGD(lr=0,05, decay=10e-6)
```

► Define loss (and performance) metric

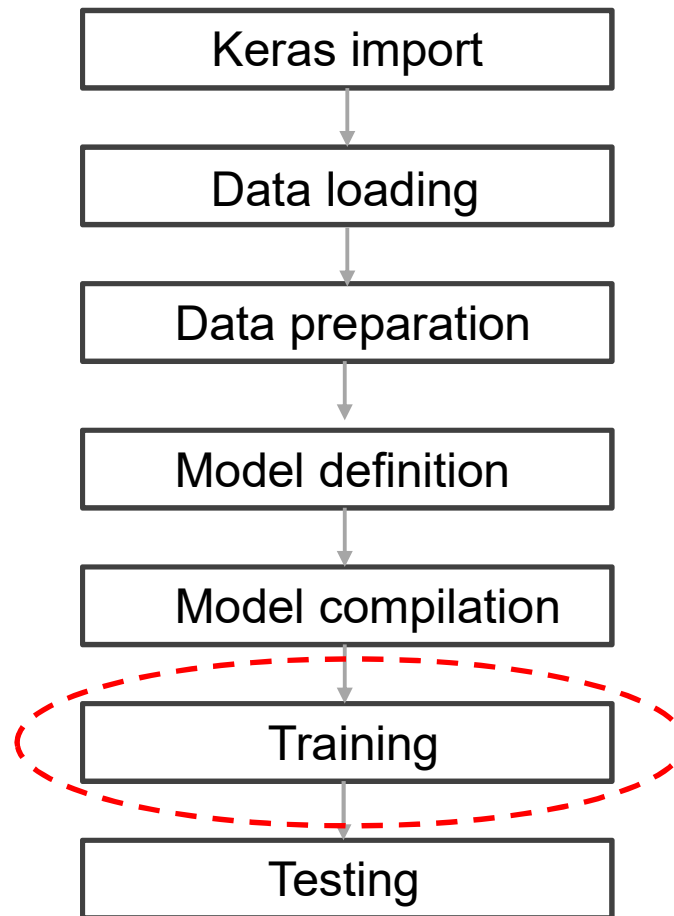
```
model.compile(optimizer=myOpt, loss='categorical_crossentropy', [...])
```

► After compiling the model, visualize it

```
model.summary()
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 2)	6
Total params: 6		

Typical Keras Dataflow



Training

► High-level `fit()` function

```
history = model.fit(x=train_data[firstSample:lastSample],
                    y=train_labels[firstSample:lastSample],
                    validation_data=(test_data, test_labels),
                    epochs=10,
                    batch_size=32,
                    shuffle=True,
                    verbose=0|1|2)
```

► Set `verbose=1` to visualize this output

```
Epoch 1/100
100/100 [=====] - 0s 2ms/step - loss: 0.9530 - acc: 0.5500
Epoch 2/100
100/100 [=====] - 0s 94us/step - loss: 0.6889 - acc: 0.5700

...

Epoch 99/100
100/100 [=====] - 0s 70us/step - loss: 0.0371 - acc: 1.0000
Epoch 100/100
100/100 [=====] - 0s 67us/step - loss: 0.0369 - acc: 1.0000
```

► Train/Validation loss/accuracy logged in history

Training with Augmentation – High Level

▶ High-level `fit_generator()` function

```
history = model.fit_generator(  
    x=train_data[firstSample:lastSample],  
    y=train_labels[firstSample:lastSample],  
    myDatagen.flow(train_data, train_labels, batch_size=32),  
    steps_per_epoch = train_images.shape[0] / 32,  
    validation_data=(test_data, test_labels),  
    epochs=10,  
    shuffle=True,  
    verbose=1,  
    workers=2  
)
```

- ▶ Also `validation_data` could be augmented in principle
 - ▶ Typically only train data augmented

Training with Augmentation – Lower Level

► First fit generator to images


```
history = model.fit_generator(  
    myDatagen.flow(train_data, train_labels, batch_size=32),  
    samples_per_epoch=len(train_data),  
    epochs=nb_epochs)
```

► Then iterate over each epoch/batch

```
for e in range(nb_epoch):  
    loss = 0  
    batchCnt = 0  
    for X_batch, Y_batch in myDatagen.flow(train_data,  
train_labels, batch_size=32):  
        loss += model.train(X_batch, Y_batch)  
        batchCnt += 1  
        if batchCnt >= len(X_train) / 32:  
            # break the loop or generator loops indefinitely  
            break  
    print ( 'Epoch' + str(e) + ' loss ' + str(loss/batchCnt))
```

Training - Monitoring

- Monitor the CPU usage via htop



Training - Monitoring

► Monitor the GPU usage via nvidia-smi

```
File Modifica Visualizza Cerca Terminale Aiuto
Every 1.0s: nvidia-smi                               sun: Mon May 27 23:10:20 2019
Mon May 27 23:10:20 2019
+-----+
| NVIDIA-SMI 410.79          Driver Version: 410.79          CUDA Version: 10.0          |
+-----+-----+
| GPU   Name           Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
+-----+-----+-----+
|  0   GeForce RTX 208...    On      | 00000000:18:00.0 Off  |          N/A         |
| 28%   27C    P8      14W / 175W | 10MiB / 10989MiB |    0%      Default   |
+-----+-----+-----+
|  1   GeForce RTX 208...    On      | 00000000:3B:00.0 Off  |          N/A         |
| 28%   28C    P8      18W / 175W | 10MiB / 10989MiB |    0%      Default   |
+-----+-----+-----+
|  2   GeForce RTX 208...    On      | 00000000:86:00.0 Off  |          N/A         |
| 39%   66C    P2     169W / 175W | 2771MiB / 10989MiB |   87%      Default   |
+-----+-----+-----+
|  3   GeForce RTX 208...    On      | 00000000:AF:00.0 Off  |          N/A         |
| 28%   33C    P8       8W / 175W | 10MiB / 10989MiB |    0%      Default   |
+-----+-----+-----+

+-----+
| Processes:                                     GPU Memory |
|  GPU       PID    Type   Process name                     Usage      |
+-----+-----+
|    2      12564    C     python                          2761MiB    |
+-----+
```

Training - Analysis

► Matlab-like plotting library

```
import matplotlib.pyplot as plt
```


► Plot the loss graph

```
plt.plot(history.history['loss'])  
plt.title('model loss')  
plt.ylabel('loss')  
plt.xlabel('iteration')  
plt.legend(['train', 'test'])  
plt.show()
```



► Plot the accuracy graph

```
plt.plot(history.history['acc'])  
plt.title('model accuracy')  
plt.ylabel('accuracy')  
plt.xlabel('epoch')  
plt.legend(['train', 'test'])  
plt.show()
```



Training – Confusion Matrix

- ▶ Another courtesy of sklearn

```
from sklearn.metrics import confusion_matrix
```

- ▶ Which class is likely confused with which

```
predictions = model.predict(test_images)
matrix = confusion_matrix(test_labels.argmax(axis=1),
                           predictions.argmax(axis=1))
print (matrix)
```

```
[[ 950   0   1   5   1   7   8   1   5   2]
 [   0 1109   2   4   0   2   3   0  15   0]
 [   7   3  943  17  18   2  11  12  12   7]
 [   3   3  14  927   1  22   2  10  22   6]
 [   1   3   4   1  891   0  35   0   4  43]
 [  12   5   5  54  14  733  16   6  34  13]
 [  14   4   1   2  15   5  913   0   4   0]
 [   1  19  38   7   7   1   0  918   4  33]
 [  11   0   5  41  11  20  11   8  853  14]
 [   7   7   9   7  30   7   2  23  13  904]]
```

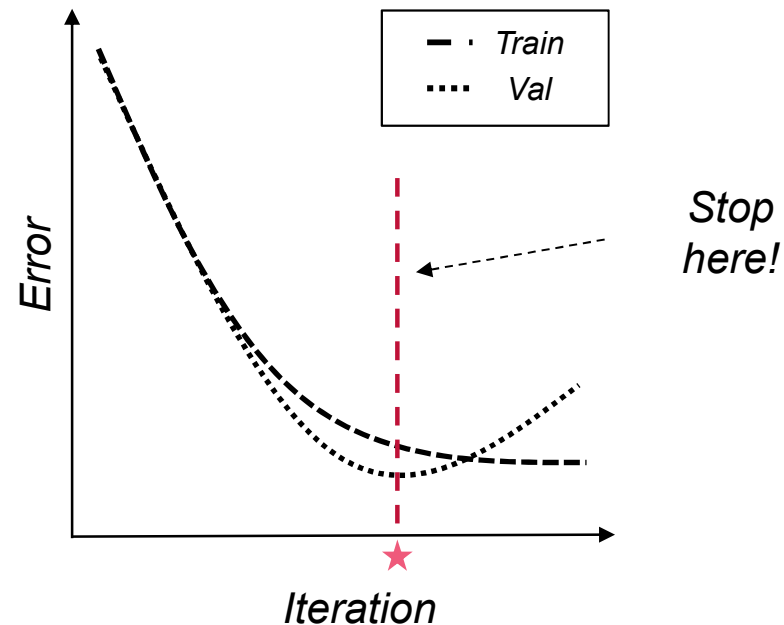
Training – Early Stop

- ▶ Useful functions for a variety of purposes

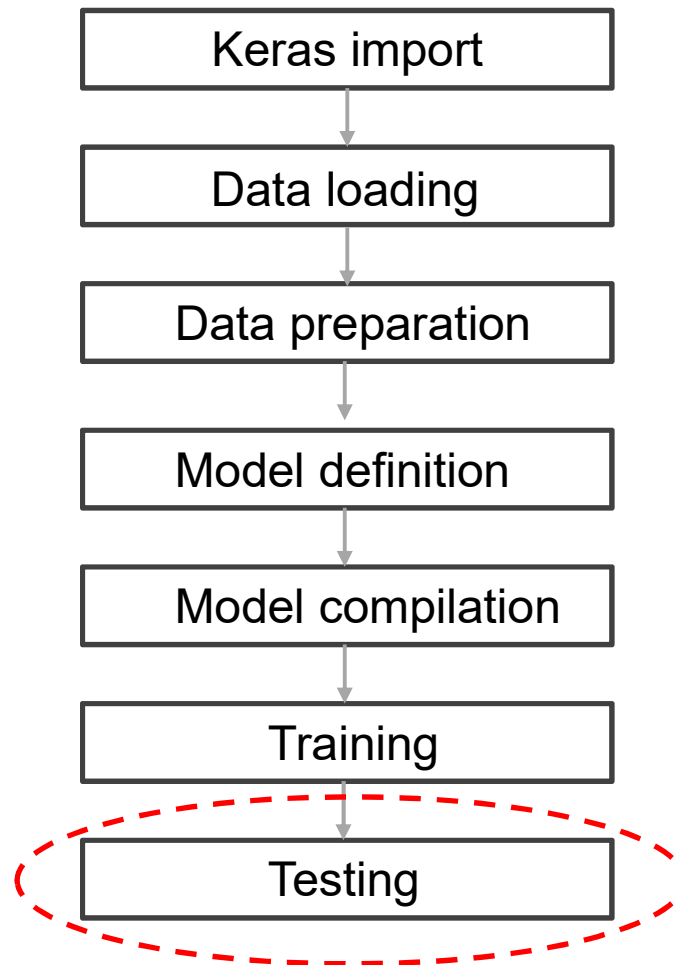
```
from keras.callbacks import EarlyStopping
```

- ▶ Ends training when validation loss stops decreasing

```
early_stopping = EarlyStopping(monitor='val_loss', patience=2)  
model.fit(..., callbacks=[early_stopping])
```



Typical Keras Dataflow



Testing

▶ Left-out samples

```
score = model.evaluate(  
    data_test[firstSample:lastSample],  
    labels_test[firstSample:lastSample]  
)
```

Model Saving / Loading

- ▶ Let us save the trained model (topology + params)

```
import os
model_name = 'trained_model.h5'
model.save(model_path)
print('Saved trained model at ' + os.getcwd() + '/' + model_path)
del(model)
```

- ▶ And let us load the trained model later on

```
from keras.models import load_model
model = load_model('trained_model.h5')
```

- ▶ Save/load only the parameters

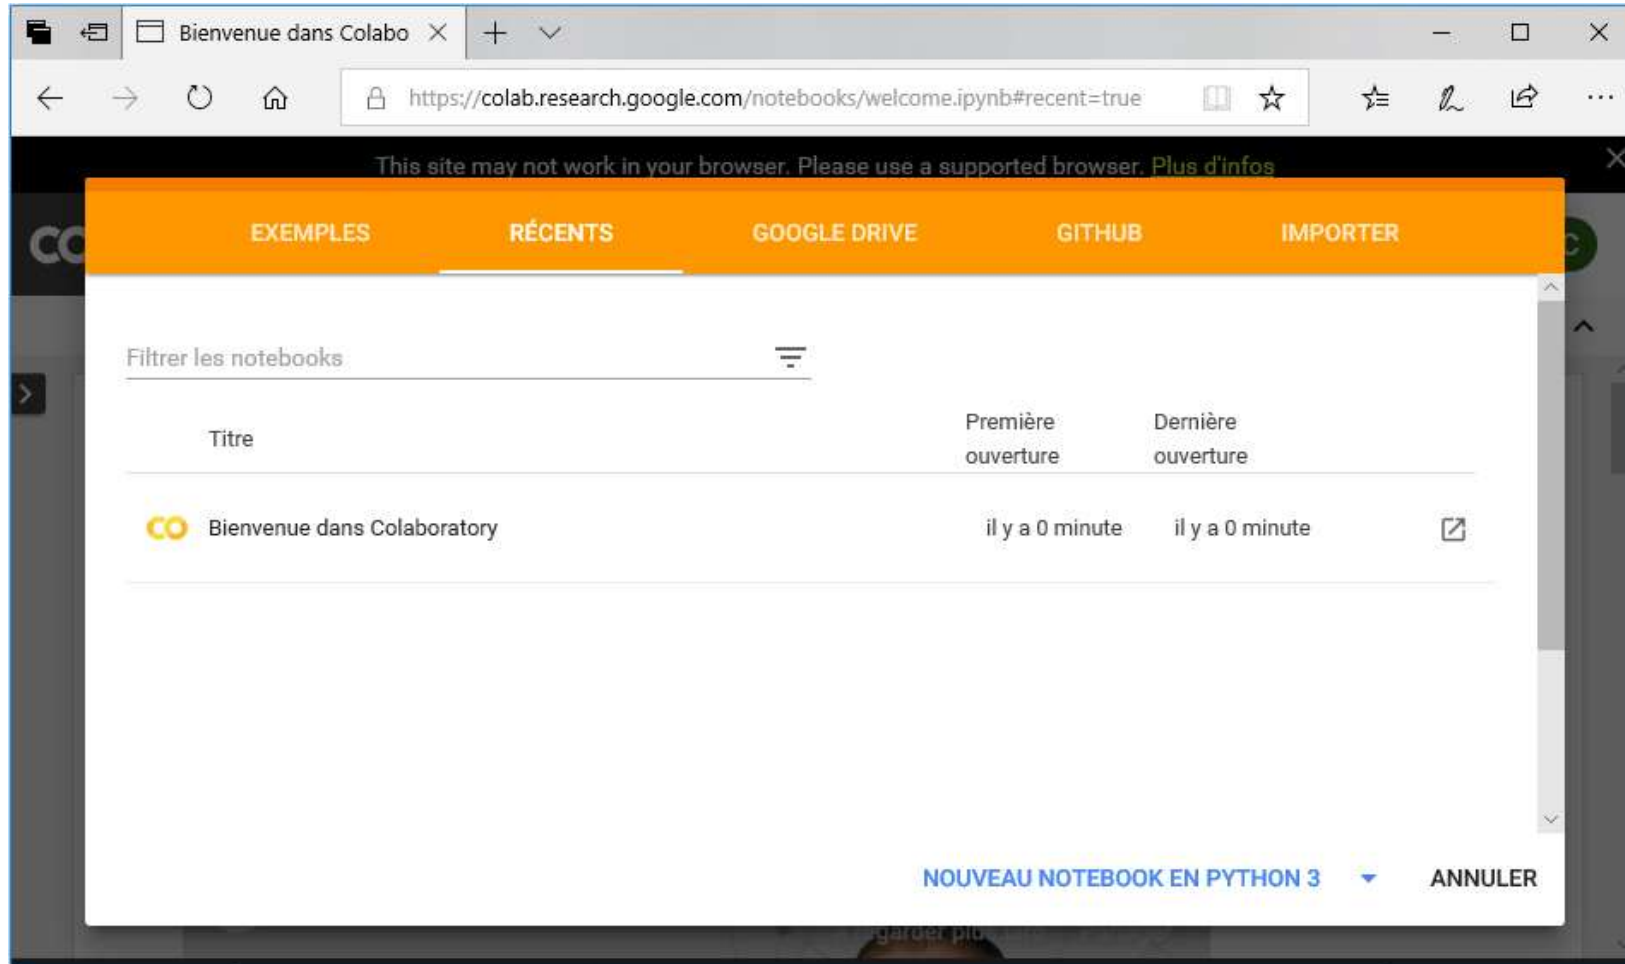
```
model.save_weights('my_model_weights.h5')

[...]

model.load_weights('my_model_weights.h5', by_name=True)
```

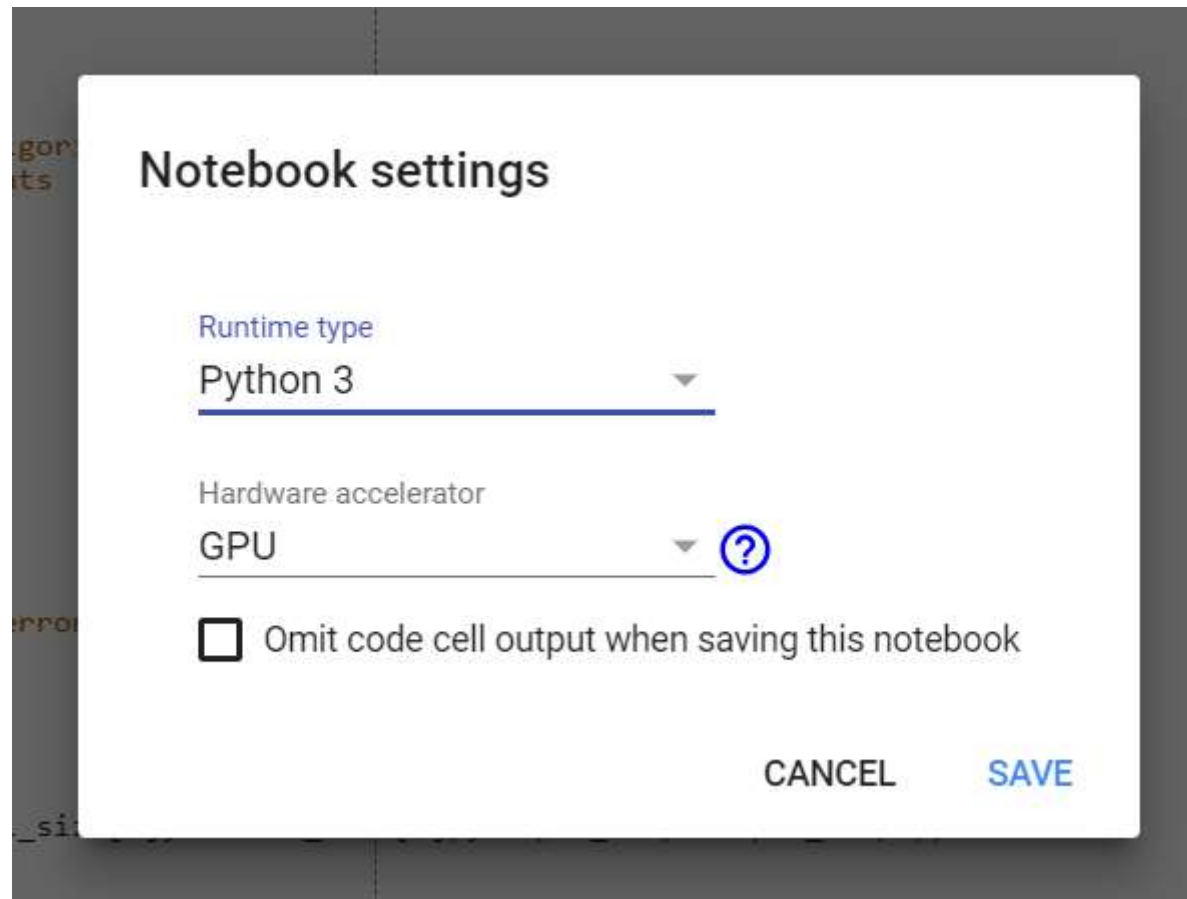
Google Colab

► Jupiter-like development environment



Google Colab

- ▶ If available, select a GPU-enabled VM («runtime»)
 - ▶ *Menu - Runtime – Change runtime type - GPU*



Training over Custom Data with Colab

- ▶ First store your own data on GDrive

- ▶ **Check privacy issues first!**

- ▶ Import the drive module from google

```
from google.colab import drive
```

- ▶ Mount your gdrive to some path

```
drive.mount('/content/drive/')
```

- ▶ Follow procedure for authentication/authorization

- ▶ Access your files as locally

```
!ls "/content/drive/My Drive/"
```

Colab Diagnostics

- ▶ Check which GPU you have available
 - ▶ «!» means system-level command (*bash*)

```
[ ] !nvidia-smi
```

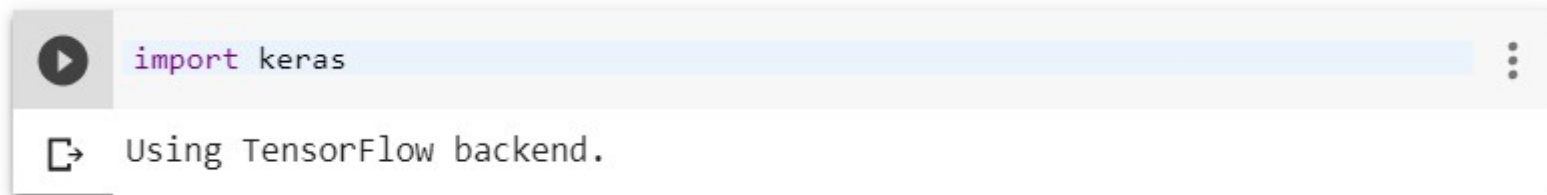
Wed Nov 21 13:29:39 2018

```
+-----+
| NVIDIA-SMI 396.44                  Driver Version: 396.44          |
+-----+-----+-----+-----+
| GPU  Name            Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp   Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
+-----+-----+-----+-----+
|   0   Tesla K80           Off      | 00000000:00:04.0 Off  |             0        |
| N/A   35C    P0       68W / 149W | 649MiB / 11441MiB |      0%      Default  |
+-----+-----+-----+-----+

+-----+
| Processes:                                             GPU Memory |
|  GPU       PID    Type    Process name                     Usage      |
+-----+-----+-----+-----+
|
```

Colab Diagnostics

- ▶ Make sure Keras uses the *TensorFlow* backend
 - ▶ «*NHWC*» data ordering required




```
import keras
```

Using TensorFlow backend.

Preliminaries

- ▶ Do not let the terminal inactive to avoid disconnection
 - ▶ Loss of all session variables and data on filesystem



Exercise

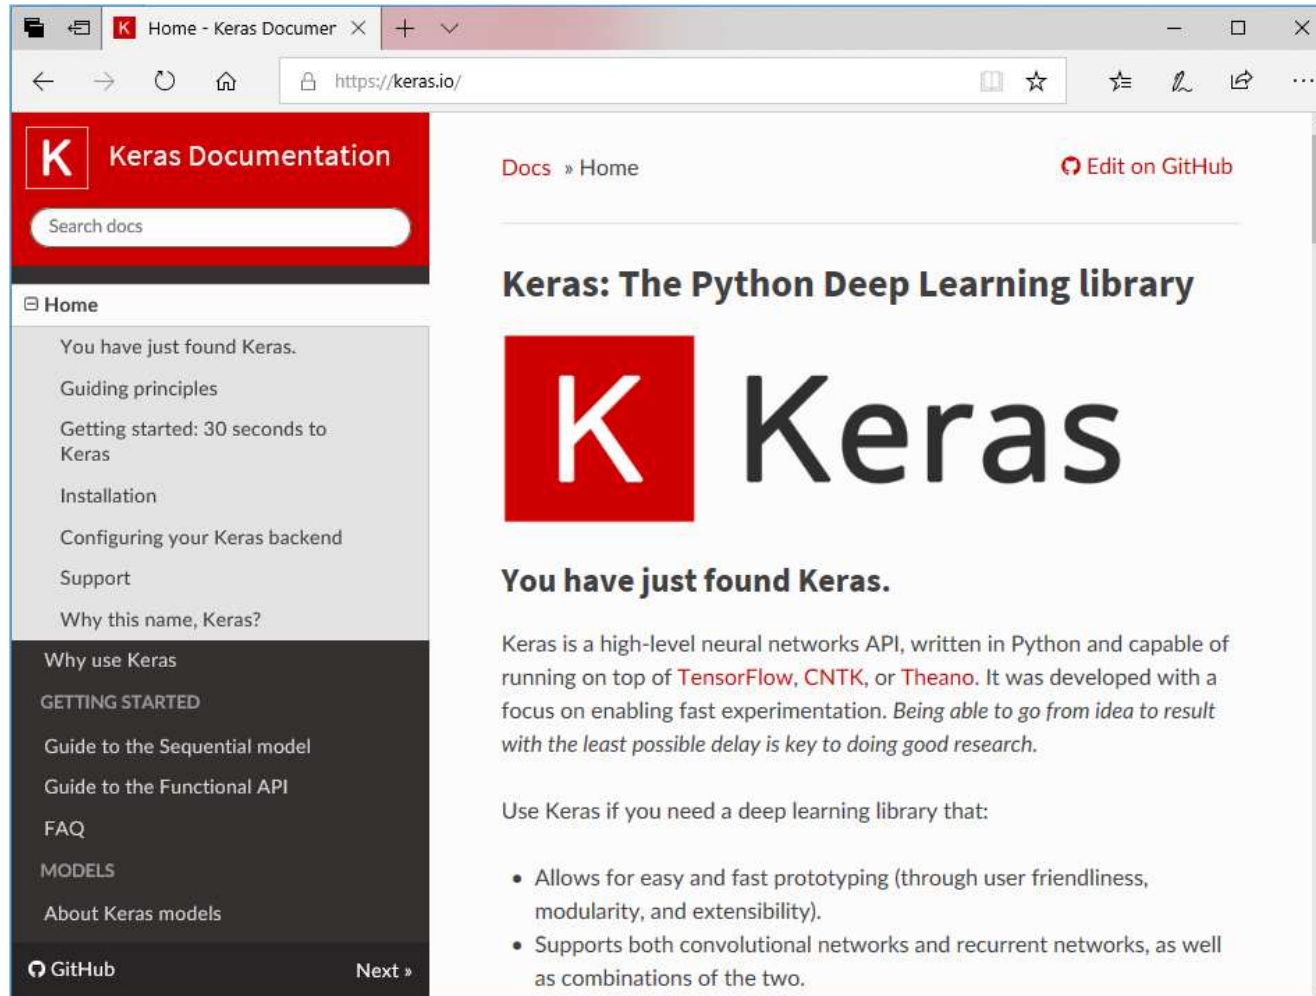


1. Generate two linearly separable classes of points and train a simple sigmoid classifier
2. Generate two non linearly separable classes of points and find the simplest FCN capable to separate them

```
data, labels = datasets.make_moons(n_samples=1000, noise=0.05, random_state=0)
```

3. Find the projection of the input data to a linearly separable space using the above trained architecture

The keras.io Website



The Keras Cheat Sheet

Python For Data Science Cheat Sheet

Keras

Learn Python for data science <https://www.DataCamp.com>



Keras

Keras is a powerful and easy-to-use deep learning library for Theano and TensorFlow that provides a high-level neural networks API to develop and evaluate deep learning models.

A Basic Example

```
>>> import numpy as np
>>> from keras.models import Sequential
>>> from keras.layers import Dense
>>> data = np.random.random((1000,20))
>>> labels = np.random.randint(2,size=(1000,1))
>>> model = Sequential()
>>> model.add(Dense(12,
>>>                 activation='relu',
>>>                 input_dim=20))
>>> model.add(Dense(1, activation='sigmoid'))
>>> model.compile(optimizer='rmsprop',
>>>               loss='binary_crossentropy',
>>>               metrics=['accuracy'])
>>> model.fit(data, labels, epochs=10, batch_size=32)
>>> predictions = model.predict(data)
```

Data

Also see NumPy, Pandas & Scikit-Learn

Your data needs to be stored as NumPy arrays or as a list of NumPy arrays. Ideally, you split the data in training and test sets, for which you can also resort to the `train_test_split` module of `sklearn.cross_validation`.

Keras Data Sets

```
>>> from keras.datasets import boston_housing,
>>>                             mnist,
>>>                             cifar10,
>>>                             imdb
>>> ix_train, y_train, ix_test, y_test = mnist.load_data()
>>> ix_train2, y_train2, ix_test2, y_test2 = boston_housing.load_data()
>>> ix_train3, y_train3, ix_test3, y_test3 = cifar10.load_data()
>>> ix_train4, y_train4, ix_test4, y_test4 = imdb.load_data(num_words=20000)
>>> num_classes = 10
```

Other

```
>>> from urllib.request import urlopen
>>> data = np.loadtxt(urlopen("https://archive.ics.uci.edu/
>>>                           ml/machine-learning-databases/plus-indians-diabetes/
>>>                           plus-indians-diabetes.data"), delimiter=","
>>> X = data[:,0:8]
>>> y = data[:,8]
```

Preprocessing

Sequence Padding

```
>>> from keras.preprocessing import sequence
>>> x_train4 = sequence.pad_sequences(x_train4, maxlen=80)
>>> x_test4 = sequence.pad_sequences(x_test4, maxlen=80)
```

One-Hot Encoding

```
>>> from keras.utils import to_categorical
>>> y_train = to_categorical(y_train, num_classes)
>>> y_test = to_categorical(y_test, num_classes)
>>> y_train3 = to_categorical(y_train3, num_classes)
>>> y_test3 = to_categorical(y_test3, num_classes)
```

Model Architecture

Sequential Model

```
>>> from keras.models import Sequential
>>> model = Sequential()
>>> model.add(Dense(12))
>>> model.add(Dense(12))
```

Multilayer Perceptron (MLP)

Binary Classification

```
>>> from keras.layers import Dense
>>> model.add(Dense(12,
>>>                 input_dim=8,
>>>                 kernel_initializer='uniform',
>>>                 activation='relu'))
>>> model.add(Dense(8, kernel_initializer='uniform', activation='relu'))
>>> model.add(Dense(1, kernel_initializer='uniform', activation='sigmoid'))
```

Multi-Class Classification

```
>>> from keras.layers import Dropout
>>> model.add(Dense(512, activation='relu', input_shape=(784,1)))
>>> model.add(Dropout(0.2))
>>> model.add(Dense(512, activation='relu'))
>>> model.add(Dropout(0.2))
>>> model.add(Dense(10, activation='softmax'))
```

Regression

```
>>> model.add(Dense(64, activation='relu', input_dim=train_data.shape[1]))
>>> model.add(Dense(1))
```

Convolutional Neural Network (CNN)

```
>>> from keras.layers import Activation, Conv2D, MaxPooling2D, Flatten
>>> model.add(Conv2D(32, (3,3), padding='same', input_shape=train_data.shape[1:]))
>>> model.add(Activation('relu'))
>>> model.add(Conv2D(32, (3,3)))
>>> model.add(Activation('relu'))
>>> model.add(MaxPooling2D(pool_size=(2,2)))
>>> model.add(Dropout(0.25))
>>> model.add(Conv2D(64, (3,3), padding='same'))
>>> model.add(Activation('relu'))
>>> model.add(Conv2D(64, (3,3)))
>>> model.add(Activation('relu'))
>>> model.add(MaxPooling2D(pool_size=(2,2)))
>>> model.add(Dropout(0.25))
>>> model.add(Flatten())
>>> model.add(Dense(512))
>>> model.add(Activation('relu'))
>>> model.add(Dropout(0.5))
>>> model.add(Dense(num_classes))
>>> model.add(Activation('softmax'))
```

Recurrent Neural Network (RNN)

```
>>> from keras.layers import Embedding, LSTM
>>> model.add(Embedding(20000,128))
>>> model.add(LSTM(128, dropout=0.2, recurrent_dropout=0.2))
>>> model.add(Dense(1, activation='sigmoid'))
```

Train and Test Sets

```
>>> from sklearn.model_selection import train_test_split
>>> x_train, x_test, y_train, y_test = train_test_split(x,
>>>                                                    y,
>>>                                                    test_size=0.33,
>>>                                                    random_state=42)
```

Standardization/Normalization

```
>>> from sklearn.preprocessing import StandardScaler
>>> scaler = StandardScaler()
>>> standardized_X = scaler.transform(x_train2)
>>> standardized_X_test = scaler.transform(x_test2)
```

Inspect Model

```
>>> model.output_shape
>>> model.summary()
>>> model.get_weights()
```

Model output shape
Model summary representation
Model configuration
List all weight tensors in the model

Compile Model

MLP: Binary Classification

```
>>> model.compile(optimizer='adam',
>>>               loss='binary_crossentropy',
>>>               metrics=['accuracy'])
```

MLP: Multi-Class Classification

```
>>> model.compile(optimizer='rmsprop',
>>>               loss='categorical_crossentropy',
>>>               metrics=['accuracy'])
```

MLP: Regression

```
>>> model.compile(optimizer='rmsprop',
>>>               loss='mse',
>>>               metrics=['mse'])
```

Recurrent Neural Network

```
>>> model.compile(loss='binary_crossentropy',
>>>               optimizer='adam',
>>>               metrics=['accuracy'])
```

Model Training

```
>>> model.fit(x_train4,
>>>          y_train4,
>>>          batch_size=32,
>>>          epochs=10,
>>>          verbose=1,
>>>          validation_data=(x_test4, y_test4))
```

Evaluate Your Model's Performance

```
>>> score = model.evaluate(x_test,
>>>                        y_test,
>>>                        batch_size=32)
```

Prediction

```
>>> model.predict(x_test4, batch_size=32)
>>> model.predict_classes(x_test4, batch_size=32)
```

Save/ Reload Models

```
>>> from keras.models import load_model
>>> model1.save('model1.h5')
>>> my_model = load_model('my_model.h5')
```

Model Fine-tuning

Optimization Parameters

```
>>> from keras.optimizers import RMSprop
>>> opt = RMSprop(lr=0.001, decay=1e-6)
>>> model.compile(loss='categorical_crossentropy',
>>>               optimizer=opt,
>>>               metrics=['accuracy'])
```

Early Stopping

```
>>> from keras.callbacks import EarlyStopping
>>> early_stopping_monitor = EarlyStopping(patience=2)
>>> model.fit(x_train4,
>>>          y_train4,
>>>          batch_size=32,
>>>          epochs=10,
>>>          validation_data=(x_test4, y_test4),
>>>          callbacks=[early_stopping_monitor])
```

DataCamp

Learn Python for Data Science interactively





Good Work!

